

Chapter 18 Interactive Notebook for Instructors

Ram Gopal, Dan Philps, and Tillman Weyde

2022

Contents

Load packages	2
Get the data and pre-process	2
Read data	2
Partition Data	2
Linear Regression Model	3
Compute functions for Residual Mean,MSE, RMSE and R_2	3
Performance of the Linear Regression Model	4
Polynomial Regression Model	4
Generate the dataset	4
Create training, validation and test sets	5
Run the Polynomial Regression	6
Evaluate the results	6
Ridge Regression	6
Basic Model	6
Hyper-parameter tuning	7
Run the Optimized Ridge Regression	7
Evaluate Performance of the optimized Ridge Regression	8
Lasso Regression	9
Basic Model	9
Hyper-parameter tuning	9
Run the Optimized Lasso Regression	9
Performance of the Optimized Lasso Regression	10
Determine non-zero coefficients	11

Neural Network	11
Large network	11
Smaller network	12
Regression Tree	12
Basic Tree	12
Tree with a maximum depth of 2	13
Harder problem	14
Data preparation	14
Build and Evaluate Models	14

Load packages

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(rattle)
```

```
## Loading required package: tibble
```

```
## Loading required package: bitops
```

```
## Rattle: A free graphical interface for data science with R.
```

```
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
```

```
## Type 'rattle()' to shake, rattle, and roll your data.
```

Get the data and pre-process

Read data

```
dataset <- read.csv('diabetes.csv', header = TRUE)
dim(dataset)
```

```
## [1] 442 11
```

Partition Data

- Note the use of `sample()` function and `sample_n()` to simplify the coding. Could be useful to explain to students the use of `apply()` functions in R.

```

set.seed(123456)
N = nrow(dataset)
cut1 = floor(0.6*N)
cut2 = floor(0.8*N)
index = sample(1:N)
train_index = index[1:cut1]
val_index = index[(cut1+1):cut2]
test_index = index[(cut2+1):N]
df_train = dataset[train_index,]
df_val = dataset[val_index,]
df_test = dataset[test_index,]
sapply(list(df_train,df_val,df_test),nrow)

```

```
## [1] 265 88 89
```

Linear Regression Model

- For this and the subsequent chapter, we will make use of the caret package that provides a uniform approach to work with a variety of predictive models. Also note the use of :: notation. This is used to indicate the package that contains the function that follows. This is helpful as often different packages may use the same function names.

```

df_train[-11] = scale(df_train[-11])
lr = caret::train(target ~ ., method='lm',data = df_train)
train_pred = predict(lr,newdata = df_train)
val_pred = predict(lr,newdata = scale(df_val[-11]))
test_pred = predict(lr,newdata = scale(df_test[-11]))

```

Compute functions for Residual Mean,MSE, RMSE and R_2

```

rm <- function(actual,pred) {
  return(mean(abs(actual-pred)))
}
mse <- function(actual,pred) {
  return(mean((pred-actual)^2))
}
rmse <- function(actual,pred) {
  return(mse(pred,actual)^0.5)
}
R_2 <- function(actual,pred){
  mean_v = rep(mean(actual),length(actual))
  SST = sum((actual-mean_v)^2)
  SSE = sum((actual-pred)^2)
  return(1-(SSE/SST))
}

```

Performance of the Linear Regression Model

```
res = data.frame()
w = rm(df_train$target,train_pred)
x = mse(df_train$target,train_pred)
y = rmse(df_train$target,train_pred)
z = R_2(df_train$target,train_pred)
res = rbind(res,c("Train",w,x,y,z))
w = rm(df_val$target,val_pred)
x = mse(df_val$target,val_pred)
y = rmse(df_val$target,val_pred)
z = R_2(df_val$target,val_pred)
res = rbind(res,c("Validation",w,x,y,z))
w = rm(df_test$target,test_pred)
x = mse(df_test$target,test_pred)
y = rmse(df_test$target,test_pred)
z = R_2(df_test$target,test_pred)
res = rbind(res,c("Test",w,x,y,z))
colnames(res) = c("Data","Residual Mean","MSE","RMSE","R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	Residual Mean	MSE	RMSE	R_2
Train	41.73	2655	51.52	0.5749
Validation	47.85	3548	59.57	0.3137
Test	43.83	3023	54.98	0.4676

Polynomial Regression Model

- Column names of the original data set.

```
colnames(df_train)
```

```
## [1] "age"    "sex"    "bmi"    "bp"     "s1"     "s2"     "s3"     "s4"
## [9] "s5"     "s6"     "target"
```

Generate the dataset

- The strategy is to Write a formula for squared and interaction terms, and use the formula to generate a new dataset that contains the polynomial terms.

```
# Degree 2 polynomial feature generation function

pf2_transform <- function(df, target_name='target') {
```

```

formula_pf2 <- as.formula(paste(target_name, '~ .^2 +',
                                paste('poly(',
                                      colnames(df)[-c(1)],
                                      ',2, raw=TRUE)[, 2]',
                                      collapse = ' + ')
                                )
                                )
output <- model.matrix(formula_pf2, data = df)
# Rewrite column names for readability
colnames_pf2 <- c("1",
                  colnames(df)[-1], # exclude target
                  paste0(colnames(df)[-1], "^2"), # include squares
                  colnames(output)[-1:(length(df)*2-1)]) # include interactions
colnames(output) <- colnames_pf2
# Convert to dataframe
output_df <- data.frame(output)
# Exclude intercept column
output_df[,1] <- NULL
return(output_df)
}

```

Create training, validation and test sets

- Create the training data set

```

train_sc_pf2 <- pf2_transform(df_train, target_name = "target")
train_sc_pf2$target = df_train$target
train_sc_pf2 = train_sc_pf2[-20]
print(colnames(train_sc_pf2))

```

```

## [1] "sex"      "bmi"      "bp"       "s1"       "s2"       "s3"       "s4"
## [8] "s5"       "s6"       "target"   "sex.2"    "bmi.2"    "bp.2"     "s1.2"
## [15] "s2.2"     "s3.2"     "s4.2"     "s5.2"     "s6.2"     "age.sex"  "age.bmi"
## [22] "age.bp"   "age.s1"   "age.s2"   "age.s3"   "age.s4"   "age.s5"   "age.s6"
## [29] "sex.bmi"  "sex.bp"   "sex.s1"   "sex.s2"   "sex.s3"   "sex.s4"   "sex.s5"
## [36] "sex.s6"   "bmi.bp"   "bmi.s1"   "bmi.s2"   "bmi.s3"   "bmi.s4"   "bmi.s5"
## [43] "bmi.s6"   "bp.s1"    "bp.s2"    "bp.s3"    "bp.s4"    "bp.s5"    "bp.s6"
## [50] "s1.s2"    "s1.s3"    "s1.s4"    "s1.s5"    "s1.s6"    "s2.s3"    "s2.s4"
## [57] "s2.s5"    "s2.s6"    "s3.s4"    "s3.s5"    "s3.s6"    "s4.s5"    "s4.s6"
## [64] "s5.s6"

```

```
dim(df_train[-10])
```

```
## [1] 265 10
```

```
dim(train_sc_pf2)
```

```
## [1] 265 64
```

- Prepare the validation and test sets

```
df_val[-11]= scale(df_val[-11])
val_sc_pf2 <- pf2_transform(df_val,target_name = "target")
val_sc_pf2$target = df_val$target
val_sc_pf2 = val_sc_pf2[-20]

df_test[-11]= scale(df_test[-11])
test_sc_pf2 <- pf2_transform(df_test,target_name = "target")
test_sc_pf2$target = df_test$target
test_sc_pf2 = test_sc_pf2[-20]
```

Run the Polynomial Regression

```
lr = caret::train(target ~ ., method='lm',data = train_sc_pf2)
train_pred = predict(lr,newdata = train_sc_pf2)
val_pred = predict(lr,newdata = val_sc_pf2)
test_pred = predict(lr,newdata = test_sc_pf2)
```

Evaluate the results

```
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	RMSE
Train	45.13
Validation	288.66
Test	231.93

Ridge Regression

Basic Model

```

ridge <- caret::train(y = train_sc_pf2$target, x = train_sc_pf2[-10],
                     method = 'glmnet',
                     tuneGrid = expand.grid(alpha = 0, lambda = 1)
                     )
train_pred = predict(ridge, newdata = train_sc_pf2)
val_pred = predict(ridge, newdata = val_sc_pf2)
test_pred = predict(ridge, newdata = test_sc_pf2)

res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[, sapply(res, is.numeric)] <- round(res[, sapply(res, is.numeric)], 4)
knitr::kable(res)

```

Data	RMSE
Train	46.67
Validation	59.98
Test	58.11

Hyper-parameter tuning

```

parameters <- c(seq(0.1, 2, by = 0.1) , seq(2, 5, 0.5) , seq(5, 25, 1))

ridge <- caret::train(y = train_sc_pf2$target,
                     x = train_sc_pf2[-10],
                     method = 'glmnet',
                     tuneGrid = expand.grid(alpha = 0, lambda = parameters) ,
                     metric = "Rsquared"
                     )
paste(" Ridge Best lambda = ", ridge$finalModel$lambdaOpt)

```

```
## [1] " Ridge Best lambda = 25"
```

Run the Optimized Ridge Regression

```

ridge_best <- caret::train(y = train_sc_pf2$target,
                          x = train_sc_pf2[-10],
                          method = 'glmnet',
                          tuneGrid = expand.grid(alpha = 0, lambda = ridge$finalModel$lambdaOpt))

```

```

train_pred = predict(ridge_best,newdata = train_sc_pf2)
val_pred = predict(ridge_best,newdata = val_sc_pf2)
test_pred = predict(ridge_best,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	48.93
Validation	57.66
Test	56.02

Evaluate Performance of the optimized Ridge Regression

```

res = data.frame()
y = R2(train_pred,df_train$target)
res = rbind(res,c("Train",y))
y = R2(val_pred,df_val$target)
res = rbind(res,c("Validation",y))
y = R2(test_pred,df_test$target)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	R_2
Train	0.6295
Validation	0.3788
Test	0.4761

Lasso Regression

Basic Model

```
lasso<-caret::train(y = train_sc_pf2$target,
  x = train_sc_pf2[-10],
  method = 'glmnet',
  tuneGrid = expand.grid(alpha = 1, lambda = 0))
train_pred = predict(lasso,newdata = train_sc_pf2)
val_pred = predict(lasso,newdata = val_sc_pf2)
test_pred = predict(lasso,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	RMSE
Train	45.51
Validation	63.96
Test	59.27

Hyper-parameter tuning

- Find the best hyper parameter lambda

```
parameters <- c(seq(0.1, 2, by =0.1) , seq(2, 5, 0.5) , seq(5, 25, 1))
lasso<-caret::train(y = train_sc_pf2$target,
  x = train_sc_pf2[-10],
  method = 'glmnet',
  tuneGrid = expand.grid(alpha = 1, lambda = parameters) ,
  metric = "Rsquared"
)
paste(" Lasso Best lambda = ",lasso$finalModel$lambdaOpt)
```

```
## [1] " Lasso Best lambda = 5"
```

Run the Optimized Lasso Regression

```

lasso_best<-caret::train(y = train_sc_pf2$target,
  x = train_sc_pf2[-10],
  method = 'glmnet',
  tuneGrid = expand.grid(alpha = 1, lambda = lasso$finalModel$lambdaOpt))
train_pred = predict(lasso_best,newdata = train_sc_pf2)
val_pred = predict(lasso_best,newdata = val_sc_pf2)
test_pred = predict(lasso_best,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	52.57
Validation	58.00
Test	53.24

Performance of the Optimized Lasso Regression

```

res = data.frame()
y = R2(train_pred,df_train$target)
res = rbind(res,c("Train",y))
y = R2(val_pred,df_val$target)
res = rbind(res,c("Validation",y))
y = R2(test_pred,df_test$target)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	R_2
Train	0.5673
Validation	0.3646
Test	0.5117

Determine non-zero coefficients

```
df = data.frame(  
  lasso = as.data.frame.matrix(coef(lasso$finalModel, lasso$finalModel$lambdaOpt))  
)  
df = subset(df, s1 > 0.1)  
df$var = row.names(df)  
knitr::kable(df[order(-df$s1), c(2, 1)] [2])
```

	s1
(Intercept)	154.3001
bp	25.3282
s6	20.6368
s1	13.4184
bmi.s6	3.9376
sex.s3	2.2168
age.s6	0.7353
bmi.bp	0.4412
age.sex	0.1728
s6.2	0.1407

Neural Network

Large network

```
mlp <- caret::train(target ~ ., data = train_sc_pf2, method='mlp', size=1000)  
  
train_pred = predict(mlp, newdata = train_sc_pf2)  
val_pred = predict(mlp, newdata = val_sc_pf2)  
test_pred = predict(mlp, newdata = test_sc_pf2)  
res = data.frame()  
y = rmse(df_train$target, train_pred)  
res = rbind(res, c("Train", y))  
y = rmse(df_val$target, val_pred)  
res = rbind(res, c("Validation", y))  
y = rmse(df_test$target, test_pred)  
res = rbind(res, c("Test", y))  
colnames(res) = c("Data", "RMSE")  
# Following converts appropriate columns to numeric type  
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})  
# Following rounds numeric values to 4 digits  
res[, sapply(res, is.numeric)] <- round(res[, sapply(res, is.numeric)], 4)  
knitr::kable(res)
```

Data	RMSE
Train	104.5
Validation	103.7

Data	RMSE
Test	107.5

Smaller network

```
mlp <- caret::train(target ~ ., data = train_sc_pf2, method='mlp', size=50)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

train_pred = predict(mlp, newdata = train_sc_pf2)
val_pred = predict(mlp, newdata = val_sc_pf2)
test_pred = predict(mlp, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[, sapply(res, is.numeric)] <- round(res[, sapply(res, is.numeric)], 4)
knitr::kable(res)
```

Data	RMSE
Train	83.34
Validation	74.52
Test	77.62

Regression Tree

Basic Tree

```
dtr <- caret::train(target ~ ., data = train_sc_pf2, method='rpart')

train_pred = predict(dtr, newdata = train_sc_pf2)
val_pred = predict(dtr, newdata = val_sc_pf2)
test_pred = predict(dtr, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
```

```

res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	60.28
Validation	66.73
Test	60.40

Tree with a maximum depth of 2

```

dtr <- caret::train(target ~ .,data = train_sc_pf2,method='rpart',
                    control = list(max_depth=2))

```

```

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

```

```

train_pred = predict(dtr,newdata = train_sc_pf2)
val_pred = predict(dtr,newdata = val_sc_pf2)
test_pred = predict(dtr,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	60.28
Validation	66.73
Test	60.40

Harder problem

Data preparation

- Read data

```
df <- read.csv("ENB2012_data.csv")
```

- Create training, validation and test datasets.

```
dataset = df[-c(11,12)]
N = nrow(dataset)
cut1 = floor(0.6*N)
cut2 = floor(0.8*N)
index = sample(1:N)
train_index = index[1:cut1]
val_index = index[(cut1+1):cut2]
test_index = index[(cut2+1):N]
df_train = dataset[train_index,]
df_val = dataset[val_index,]
df_test = dataset[test_index,]
sapply(list(df_train,df_val,df_test),nrow)
```

```
## [1] 460 154 154
```

```
df_train[-10] = scale(df_train[-10])
```

Build and Evaluate Models

```
res = data.frame()
# Linear Regression
lr = caret::train(Y2 ~ ., method='lm', data = df_train)
train_pred = predict(lr, newdata = df_train)
val_pred = predict(lr, newdata = scale(df_val[-10]))
test_pred = predict(lr, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Linear Regression", y1, y2, y3))
# Ridge Regression
ridge <- caret::train(y = df_train$Y2, x = df_train[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 0, lambda = 1)
)
train_pred = predict(ridge, newdata = df_train)
val_pred = predict(ridge, newdata = scale(df_val[-10]))
test_pred = predict(ridge, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
```

```

y3 = rmse(df_test$Y2,test_pred)
res = rbind(res,c("Ridge Regression",y1,y2,y3))
# Lasso Regression
lasso <- caret::train(y = df_train$Y2,x = df_train[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 1, lambda = 0)
                      )
train_pred = predict(lasso,newdata = df_train)
val_pred = predict(lasso,newdata = scale(df_val[-10]))
test_pred = predict(lasso,newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2,train_pred)
y2 = rmse(df_val$Y2,val_pred)
y3 = rmse(df_test$Y2,test_pred)
res = rbind(res,c("Lasso Regression",y1,y2,y3))
# Neural Net
nn <- caret::train(y = df_train$Y2,x = df_train[-10],
                   method = 'mlp')
train_pred = predict(nn,newdata = df_train)
val_pred = predict(nn,newdata = scale(df_val[-10]))
test_pred = predict(nn,newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2,train_pred)
y2 = rmse(df_val$Y2,val_pred)
y3 = rmse(df_test$Y2,test_pred)
res = rbind(res,c("Neural Net",y1,y2,y3))
# Regression Tree
dt <- caret::train(y = df_train$Y2,x = df_train[-10],
                   method = 'rpart')
train_pred = predict(dt,newdata = df_train)
val_pred = predict(dt,newdata = scale(df_val[-10]))
test_pred = predict(dt,newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2,train_pred)
y2 = rmse(df_val$Y2,val_pred)
y3 = rmse(df_test$Y2,test_pred)
res = rbind(res,c("Regression Tree",y1,y2,y3))

colnames(res) = c("Model","Train","Validation","Test")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Model	Train	Validation	Test
Linear Regression	1.901	2.014	2.026
Ridge Regression	2.223	2.453	2.481
Lasso Regression	1.903	2.015	2.022
Neural Net	2.506	2.529	2.574
Regression Tree	2.989	3.219	3.325