

Chapter 18 Interactive Notebook for Instructors

Ram Gopal, Dan Philps, and Tillman Weyde

2022

Contents

Load packages	2
Get the data and pre-process	2
Read data	2
Partition Data	2
Linear Regression Model	3
Compute functions for Residual Mean,MSE, RMSE and R_2	4
Performance of the Linear Regression Model	4
Polynomial Regression Model	5
Generate the dataset	5
Create training, validation and test sets	5
Run the Polynomial Regression	6
Evaluate the results	6
Ridge Regression	7
Basic Model	7
Hyper-parameter tuning	8
Run the Optimized Ridge Regression	8
Evaluate Performance of the optimized Ridge Regression	9
Lasso Regression	9
Basic Model	9
Hyper-parameter tuning	10
Run the Optimized Lasso Regression	10
Performance of the Optimized Lasso Regression	11
Determine non-zero coefficients	11

Neural Network	11
Large network	11
Smaller network	12
Regression Tree	13
Basic Tree	13
Tree with a maximum depth of 2	13
Harder problem	14
Data preparation	14
Build and Evaluate Models	14

Load packages

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(rattle)
```

```
## Loading required package: tibble
```

```
## Loading required package: bitops
```

```
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

Get the data and pre-process

Read data

```
dataset <- read.csv('../data/diabetes.csv', header = TRUE)
dim(dataset)
```

```
## [1] 442 11
```

Partition Data

- Note the use of `sample()` function and `apply()` to simplify the coding. Could be useful to explain to students the use of `apply()` functions in R.

```

set.seed(123456)
N = nrow(dataset)
cut1 = floor(0.6*N)
cut2 = floor(0.8*N)
index = sample(1:N)
train_index = index[1:cut1]
val_index = index[(cut1+1):cut2]
test_index = index[(cut2+1):N]
df_train = dataset[train_index,]
df_val = dataset[val_index,]
df_test = dataset[test_index,]
sapply(list(df_train,df_val,df_test),nrow)

```

```
## [1] 265 88 89
```

Linear Regression Model

- For this and the subsequent chapter, we will make use of the caret package that provides a uniform approach to work with a variety of predictive models. Also note the use of :: notation. This is used to indicate the package that contains the function that follows. This is helpful as often different packages may use the same function names.

```

means = apply(df_train[-11],2,mean)
sds = apply(df_train[-11],2,sd)
scalefun = function(x){
  return((x-means)/sds)
}
df_train[-11] = data.frame(sapply(df_train[-11],scalefun))
df_val[-11] = data.frame(sapply(df_val[-11],scalefun))
df_test[-11] = data.frame(sapply(df_test[-11],scalefun))
printfun = function(x){
  means = apply(x[-11],2,mean)
  sds = apply(x[-11],2,sd)
  print(paste(mean(means),mean(sds)))
}
printfun(df_train)

```

```
## [1] "0.000554657253475731 1.00044597757519"
```

```
printfun(df_val)
```

```
## [1] "-0.0182503268108164 0.980026441948433"
```

```
printfun(df_test)
```

```
## [1] "0.0533555353079712 0.987301726800513"
```

```
lr = caret::train(target ~ ., method='lm', data = df_train)
train_pred = predict(lr, newdata = df_train)
val_pred = predict(lr, newdata = scale(df_val[-11]))
test_pred = predict(lr, newdata = scale(df_test[-11]))
```

Compute functions for Residual Mean, MSE, RMSE and R₂

```
rm <- function(actual, pred) {
  return(mean(abs(actual-pred)))
}
mse <- function(actual, pred) {
  return(mean((pred-actual)^2))
}
rmse <- function(actual, pred) {
  return(mse(pred, actual)^0.5)
}
R_2 <- function(actual, pred){
  mean_v = rep(mean(actual), length(actual))
  SST = sum((actual-mean_v)^2)
  SSE = sum((actual-pred)^2)
  return(1-(SSE/SST))
}
```

Performance of the Linear Regression Model

```
res = data.frame()
w = rm(df_train$target, train_pred)
x = mse(df_train$target, train_pred)
y = rmse(df_train$target, train_pred)
z = R_2(df_train$target, train_pred)
res = rbind(res, c("Train", w, x, y, z))
w = rm(df_val$target, val_pred)
x = mse(df_val$target, val_pred)
y = rmse(df_val$target, val_pred)
z = R_2(df_val$target, val_pred)
res = rbind(res, c("Validation", w, x, y, z))
w = rm(df_test$target, test_pred)
x = mse(df_test$target, test_pred)
y = rmse(df_test$target, test_pred)
z = R_2(df_test$target, test_pred)
res = rbind(res, c("Test", w, x, y, z))
colnames(res) = c("Data", "Residual Mean", "MSE", "RMSE", "R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[, sapply(res, is.numeric)] <- round(res[, sapply(res, is.numeric)], 4)
knitr::kable(res)
```

Data	Residual Mean	MSE	RMSE	R_2
Train	41.79	2659	51.56	0.5742
Validation	47.58	3497	59.13	0.3238
Test	43.69	3014	54.90	0.4693

Polynomial Regression Model

- Column names of the original data set.

```
colnames(df_train)
```

```
## [1] "age"    "sex"    "bmi"    "bp"     "s1"     "s2"     "s3"     "s4"
## [9] "s5"     "s6"     "target"
```

Generate the dataset

- The strategy is to Write a formula for squared and interaction terms, and use the formula to generate a new dataset that contains the polynomial terms.

```
# Degree 2 polynomial feature generation function

pf2_transform <- function(df, target_name='target') {
  formula_pf2 <- as.formula(paste(target_name, '~ .^2 +',
                                   paste('poly(',
                                           colnames(df)[-c(1)],
                                           ',2, raw=TRUE)[, 2]',
                                           collapse = ' + ')
                               )
  )
  output <- model.matrix(formula_pf2, data = df)
  # Rewrite column names for readability
  colnames_pf2 <- c("1",
                    colnames(df)[-1],           # exclude target
                    paste0(colnames(df)[-1], "^2"), # include squares
                    colnames(output)[-1:(length(df)*2-1)]) # include interactions
  colnames(output) <- colnames_pf2
  # Convert to dataframe
  output_df <- data.frame(output)
  # Exclude intercept column
  output_df[,1] <- NULL
  return(output_df)
}
```

Create training, validation and test sets

- Create the training data set

```
train_sc_pf2 <- pf2_transform(df_train,target_name = "target")
train_sc_pf2$target= df_train$target
train_sc_pf2 = train_sc_pf2[-20]
print(colnames(train_sc_pf2))
```

```
## [1] "sex"      "bmi"      "bp"       "s1"       "s2"       "s3"       "s4"
## [8] "s5"       "s6"       "target"   "sex.2"    "bmi.2"    "bp.2"     "s1.2"
## [15] "s2.2"     "s3.2"     "s4.2"     "s5.2"     "s6.2"     "age.sex"  "age.bmi"
## [22] "age.bp"   "age.s1"   "age.s2"   "age.s3"   "age.s4"   "age.s5"   "age.s6"
## [29] "sex.bmi"  "sex.bp"   "sex.s1"   "sex.s2"   "sex.s3"   "sex.s4"   "sex.s5"
## [36] "sex.s6"   "bmi.bp"   "bmi.s1"   "bmi.s2"   "bmi.s3"   "bmi.s4"   "bmi.s5"
## [43] "bmi.s6"   "bp.s1"    "bp.s2"    "bp.s3"    "bp.s4"    "bp.s5"    "bp.s6"
## [50] "s1.s2"    "s1.s3"    "s1.s4"    "s1.s5"    "s1.s6"    "s2.s3"    "s2.s4"
## [57] "s2.s5"    "s2.s6"    "s3.s4"    "s3.s5"    "s3.s6"    "s4.s5"    "s4.s6"
## [64] "s5.s6"
```

```
dim(df_train[-10])
```

```
## [1] 265 10
```

```
dim(train_sc_pf2)
```

```
## [1] 265 64
```

- Prepare the validation and test sets

```
df_val[-11]= scale(df_val[-11])
val_sc_pf2 <- pf2_transform(df_val,target_name = "target")
val_sc_pf2$target = df_val$target
val_sc_pf2 = val_sc_pf2[-20]
```

```
df_test[-11]= scale(df_test[-11])
test_sc_pf2 <- pf2_transform(df_test,target_name = "target")
test_sc_pf2$target = df_test$target
test_sc_pf2 = test_sc_pf2[-20]
```

Run the Polynomial Regression

```
lr = caret::train(target ~ ., method='lm',data = train_sc_pf2)
train_pred = predict(lr,newdata = train_sc_pf2)
val_pred = predict(lr,newdata = val_sc_pf2)
test_pred = predict(lr,newdata = test_sc_pf2)
```

Evaluate the results

```

res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <- round(res[,sapply(res, is.numeric)], 4)
knitr::kable(res)

```

Data	RMSE
Train	44.95
Validation	71.05
Test	77.25

Ridge Regression

Basic Model

```

ridge <- caret::train(y = train_sc_pf2$target, x = train_sc_pf2[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 0, lambda = 1)
                    )
train_pred = predict(ridge, newdata = train_sc_pf2)
val_pred = predict(ridge, newdata = val_sc_pf2)
test_pred = predict(ridge, newdata = test_sc_pf2)

res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <- round(res[,sapply(res, is.numeric)], 4)
knitr::kable(res)

```

Data	RMSE
Train	46.66

Data	RMSE
Validation	59.19
Test	57.10

Hyper-parameter tuning

```
parameters <- c(seq(0.1, 2, by =0.1) , seq(2, 5, 0.5) , seq(5, 25, 1))

ridge<-caret::train(y = train_sc_pf2$target,
                    x = train_sc_pf2[-10],
                    method = 'glmnet',
                    tuneGrid = expand.grid(alpha = 0, lambda = parameters) ,
                    metric = "Rsquared"
                    )
paste(" Ridge Best lambda = ",ridge$finalModel$lambdaOpt)
```

```
## [1] " Ridge Best lambda = 25"
```

Run the Optimized Ridge Regression

```
ridge_best<-caret::train(y = train_sc_pf2$target,
                        x = train_sc_pf2[-10],
                        method = 'glmnet',
                        tuneGrid = expand.grid(alpha = 0, lambda = ridge$finalModel$lambdaOpt))
train_pred = predict(ridge_best,newdata = train_sc_pf2)
val_pred = predict(ridge_best,newdata = val_sc_pf2)
test_pred = predict(ridge_best,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	RMSE
Train	49.03
Validation	57.48
Test	55.76

Evaluate Performance of the optimized Ridge Regression

```
res = data.frame()
y = R2(train_pred,df_train$target)
res = rbind(res,c("Train",y))
y = R2(val_pred,df_val$target)
res = rbind(res,c("Validation",y))
y = R2(test_pred,df_test$target)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	R_2
Train	0.6275
Validation	0.3794
Test	0.4754

Lasso Regression

Basic Model

```
lasso<-caret::train(y = train_sc_pf2$target,
  x = train_sc_pf2[-10],
  method = 'glmnet',
  tuneGrid = expand.grid(alpha = 1, lambda = 0))
train_pred = predict(lasso,newdata = train_sc_pf2)
val_pred = predict(lasso,newdata = val_sc_pf2)
test_pred = predict(lasso,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	RMSE
Train	45.42
Validation	62.98
Test	65.22

Hyper-parameter tuning

- Find the best hyper parameter lambda

```
parameters <- c(seq(0.1, 2, by =0.1) , seq(2, 5, 0.5) , seq(5, 25, 1))
lasso<-caret::train(y = train_sc_pf2$target,
                    x = train_sc_pf2[-10],
                    method = 'glmnet',
                    tuneGrid = expand.grid(alpha = 1, lambda = parameters) ,
                    metric = "Rsquared"
                    )
paste(" Lasso Best lambda = ",lasso$finalModel$lambdaOpt)
```

```
## [1] " Lasso Best lambda = 5"
```

Run the Optimized Lasso Regression

```
lasso_best<-caret::train(y = train_sc_pf2$target,
                        x = train_sc_pf2[-10],
                        method = 'glmnet',
                        tuneGrid = expand.grid(alpha = 1, lambda = lasso$finalModel$lambdaOpt))
train_pred = predict(lasso_best,newdata = train_sc_pf2)
val_pred = predict(lasso_best,newdata = val_sc_pf2)
test_pred = predict(lasso_best,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	RMSE
Train	52.63
Validation	57.74
Test	53.41

Performance of the Optimized Lasso Regression

```
res = data.frame()
y = R2(train_pred,df_train$target)
res = rbind(res,c("Train",y))
y = R2(val_pred,df_val$target)
res = rbind(res,c("Validation",y))
y = R2(test_pred,df_test$target)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	R_2
Train	0.5663
Validation	0.3677
Test	0.5078

Determine non-zero coefficients

```
df = data.frame(
  lasso = as.data.frame.matrix(coef(lasso$finalModel, lasso$finalModel$lambdaOpt))
)
df =subset(df, s1>0.1)
df$var = row.names(df)
knitr::kable(df[order(-df$s1),c(2,1)][2])
```

	s1
(Intercept)	153.1769
bp	25.3203
s6	19.9792
s1	13.1870
bmi.s6	3.7814
sex.s3	2.0960
age.s6	1.1315
s6.2	0.2655
bmi.bp	0.2125

Neural Network

Large network

```

mlp <- caret::train(target ~ .,data = train_sc_pf2,method='mlp',size=1000)

train_pred = predict(mlp,newdata = train_sc_pf2)
val_pred = predict(mlp,newdata = val_sc_pf2)
test_pred = predict(mlp,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	102.7
Validation	100.9
Test	107.1

Smaller network

```

mlp <- caret::train(target ~ .,data = train_sc_pf2,method='mlp',size=50)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

train_pred = predict(mlp,newdata = train_sc_pf2)
val_pred = predict(mlp,newdata = val_sc_pf2)
test_pred = predict(mlp,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	86.42
Validation	83.83
Test	86.49

Regression Tree

Basic Tree

```
dtr <- caret::train(target ~ ., data = train_sc_pf2, method='rpart')

train_pred = predict(dtr, newdata = train_sc_pf2)
val_pred = predict(dtr, newdata = val_sc_pf2)
test_pred = predict(dtr, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <- round(res[,sapply(res, is.numeric)], 4)
knitr::kable(res)
```

Data	RMSE
Train	59.71
Validation	66.82
Test	60.15

Tree with a maximum depth of 2

```
dtr <- caret::train(target ~ ., data = train_sc_pf2, method='rpart',
  control = list(max_depth=2))
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.
```

```
train_pred = predict(dtr, newdata = train_sc_pf2)
val_pred = predict(dtr, newdata = val_sc_pf2)
test_pred = predict(dtr, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
```

```

res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	59.71
Validation	66.82
Test	60.15

Harder problem

Data preparation

- Read data

```
df <- read.csv("../data/ENB2012_data.csv")
```

- Create training, validation and test datasets.

```

dataset = df[-c(11,12)]
N = nrow(dataset)
cut1 = floor(0.6*N)
cut2 = floor(0.8*N)
index = sample(1:N)
train_index = index[1:cut1]
val_index = index[(cut1+1):cut2]
test_index = index[(cut2+1):N]
df_train = dataset[train_index,]
df_val = dataset[val_index,]
df_test = dataset[test_index,]
sapply(list(df_train,df_val,df_test),nrow)

```

```
## [1] 460 154 154
```

```
df_train[-10] = scale(df_train[-10])
```

Build and Evaluate Models

```

res = data.frame()
# Linear Regression
lr = caret::train(Y2 ~ ., method='lm', data = df_train)
train_pred = predict(lr, newdata = df_train)
val_pred = predict(lr, newdata = scale(df_val[-10]))
test_pred = predict(lr, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Linear Regression", y1, y2, y3))
# Ridge Regression
ridge <- caret::train(y = df_train$Y2, x = df_train[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 0, lambda = 1)
)
train_pred = predict(ridge, newdata = df_train)
val_pred = predict(ridge, newdata = scale(df_val[-10]))
test_pred = predict(ridge, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Ridge Regression", y1, y2, y3))
# Lasso Regression
lasso <- caret::train(y = df_train$Y2, x = df_train[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 1, lambda = 0)
)
train_pred = predict(lasso, newdata = df_train)
val_pred = predict(lasso, newdata = scale(df_val[-10]))
test_pred = predict(lasso, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Lasso Regression", y1, y2, y3))
# Neural Net
nn <- caret::train(y = df_train$Y2, x = df_train[-10],
                   method = 'mlp')
train_pred = predict(nn, newdata = df_train)
val_pred = predict(nn, newdata = scale(df_val[-10]))
test_pred = predict(nn, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Neural Net", y1, y2, y3))
# Regression Tree
dt <- caret::train(y = df_train$Y2, x = df_train[-10],
                   method = 'rpart')
train_pred = predict(dt, newdata = df_train)
val_pred = predict(dt, newdata = scale(df_val[-10]))
test_pred = predict(dt, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)

```

```

res = rbind(res,c("Regression Tree",y1,y2,y3))

colnames(res) = c("Model","Train","Validation","Test")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Model	Train	Validation	Test
Linear Regression	1.921	2.140	1.992
Ridge Regression	2.308	2.241	2.540
Lasso Regression	1.923	2.139	1.993
Neural Net	3.228	3.547	3.324
Regression Tree	3.006	3.101	3.301