

# Chapter 18: Regression Models in Machine Learning

Ram Gopal, Dan Philps, and Tillman Weyde

2022

## Contents

<b>Load packages</b>	<b>1</b>
<b>Get the data and pre-process</b>	<b>2</b>
Read data . . . . .	2
Partition Data . . . . .	2
<b>Regression Model and predict</b>	<b>3</b>
<b>Compute functions for Residual Mean,MSE, RMSE and R_2</b>	<b>3</b>
<b>Polynomial</b>	<b>4</b>
Create data sets with polynomials for training, validation and test sets. . . . .	5
<b>Parameter Shrinkage with Ridge Regression</b>	<b>6</b>
<b>Run best ridge regression</b>	<b>7</b>
<b>Lasso Regression</b>	<b>8</b>
<b>Run best Lasso regression</b>	<b>9</b>
Lasso non zero coefficients . . . . .	10
<b>Neural Network</b>	<b>11</b>
Regression Trees . . . . .	12
<b>Harder problem</b>	<b>13</b>

## Load packages

```
library(caret)
```

```
## Loading required package: ggplot2

## Loading required package: lattice

library(rattle)

## Loading required package: tibble

## Loading required package: bitops

## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

## Get the data and pre-process

### Read data

```
dataset <- read.csv('../data/diabetes.csv', header = TRUE)
dim(dataset)
```

```
## [1] 442  11
```

### Partition Data

```
set.seed(123456)
N = nrow(dataset)
cut1 = floor(0.6*N)
cut2 = floor(0.8*N)
index = sample(1:N)
train_index = index[1:cut1]
val_index = index[(cut1+1):cut2]
test_index = index[(cut2+1):N]
df_train = dataset[train_index,]
df_val = dataset[val_index,]
df_test = dataset[test_index,]
sapply(list(df_train,df_val,df_test),nrow)
```

```
## [1] 265  88  89
```

```
means = apply(df_train[-11],2,mean)
sds = apply(df_train[-11],2,sd)
scalefun = function(x){
  return((x-means)/sds)
}
df_train[-11] = data.frame(sapply(df_train[-11],scalefun))
df_val[-11] = data.frame(sapply(df_val[-11],scalefun))
df_test[-11] = data.frame(sapply(df_test[-11],scalefun))
```

```

printfun = function(x){
  means = apply(x[-11],2,mean)
  sds = apply(x[-11],2,sd)
  print(paste(mean(means),mean(sds)))
}
printfun(df_train)

```

```
## [1] "0.000554657253475731 1.00044597757519"
```

```
printfun(df_val)
```

```
## [1] "-0.0182503268108164 0.980026441948433"
```

```
printfun(df_test)
```

```
## [1] "0.0533555353079712 0.987301726800513"
```

## Regression Model and predict

```

lr = caret::train(target ~ ., method='lm',data = df_train)
train_pred = predict(lr,newdata = df_train)
val_pred = predict(lr,newdata = df_val[-11])
test_pred = predict(lr,newdata = df_test[-11])

```

## Compute functions for Residual Mean,MSE, RMSE and R\_2

```

rm <- function(actual,pred) {
  return(mean(abs(actual-pred)))
}
mse <- function(actual,pred) {
  return(mean((pred-actual)^2))
}
rmse <- function(actual,pred) {
  return(mse(pred,actual)^0.5)
}
R_2 <- function(actual,pred){
  mean_v = rep(mean(actual),length(actual))
  SST = sum((actual-mean_v)^2)
  SSE = sum((actual-pred)^2)
  return(1-(SSE/SST))
}

```

```

res = data.frame()
w = rm(df_train$target,train_pred)
x = mse(df_train$target,train_pred)

```

```

y = rmse(df_train$target,train_pred)
z = R_2(df_train$target,train_pred)
res = rbind(res,c("Train",w,x,y,z))
w = rm(df_val$target,val_pred)
x = mse(df_val$target,val_pred)
y = rmse(df_val$target,val_pred)
z = R_2(df_val$target,val_pred)
res = rbind(res,c("Validation",w,x,y,z))
w = rm(df_test$target,test_pred)
x = mse(df_test$target,test_pred)
y = rmse(df_test$target,test_pred)
z = R_2(df_test$target,test_pred)
res = rbind(res,c("Test",w,x,y,z))
colnames(res) = c("Data","Residual Mean","MSE","RMSE","R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	Residual Mean	MSE	RMSE	R_2
Train	41.79	2659	51.56	0.5742
Validation	47.68	3467	58.88	0.3294
Test	44.06	3028	55.03	0.4668

## Polynomial

- Column names

```
colnames(df_train)
```

```
## [1] "age"    "sex"    "bmi"    "bp"     "s1"     "s2"     "s3"     "s4"
## [9] "s5"     "s6"     "target"
```

- Write formula for intercept, raw features, squared features, and interactions

```

# Degree 2 polynomial feature generation function

pf2_transform <- function(df, target_name='target') {
  formula_pf2 <- as.formula(paste(target_name, '~ .^2 +',
                                   paste('poly(',
                                           colnames(df)[-c(1)],
                                           ',2, raw=TRUE)[, 2]',
                                           collapse = ' + ')
                               )
  )

  output <- model.matrix(formula_pf2, data = df)
  # Rewrite column names for readability
  colnames_pf2 <- c("1",

```

```

        colnames(df)[-1],          # exclude target
        paste0(colnames(df)[-1], "^2"), # include squares
        colnames(output)[-1:(length(df)*2-1)]) # include interactions
    colnames(output) <- colnames_pf2
    # Convert to dataframe
    output_df <- data.frame(output)
    # Exclude intercept column
    output_df[,1] <- NULL
    return(output_df)
}

```

Create data sets with polynomials for training, validation and test sets.

- Create the training data set

```

train_sc_pf2 <- pf2_transform(df_train, target_name = "target")
train_sc_pf2$target = df_train$target
train_sc_pf2 = train_sc_pf2[-20]
print(colnames(train_sc_pf2))

```

```

## [1] "sex"      "bmi"      "bp"       "s1"       "s2"       "s3"       "s4"
## [8] "s5"       "s6"       "target"   "sex.2"    "bmi.2"    "bp.2"     "s1.2"
## [15] "s2.2"     "s3.2"     "s4.2"     "s5.2"     "s6.2"     "age.sex"  "age.bmi"
## [22] "age.bp"   "age.s1"   "age.s2"   "age.s3"   "age.s4"   "age.s5"   "age.s6"
## [29] "sex.bmi"  "sex.bp"   "sex.s1"   "sex.s2"   "sex.s3"   "sex.s4"   "sex.s5"
## [36] "sex.s6"   "bmi.bp"   "bmi.s1"   "bmi.s2"   "bmi.s3"   "bmi.s4"   "bmi.s5"
## [43] "bmi.s6"   "bp.s1"    "bp.s2"    "bp.s3"    "bp.s4"    "bp.s5"    "bp.s6"
## [50] "s1.s2"    "s1.s3"    "s1.s4"    "s1.s5"    "s1.s6"    "s2.s3"    "s2.s4"
## [57] "s2.s5"    "s2.s6"    "s3.s4"    "s3.s5"    "s3.s6"    "s4.s5"    "s4.s6"
## [64] "s5.s6"

```

```
dim(df_train[-10])
```

```
## [1] 265 10
```

```
dim(train_sc_pf2)
```

```
## [1] 265 64
```

- Prepare the validation and test sets

```

df_val[-11] = scale(df_val[-11])
val_sc_pf2 <- pf2_transform(df_val, target_name = "target")
val_sc_pf2$target = df_val$target
val_sc_pf2 = val_sc_pf2[-20]

df_test[-11] = scale(df_test[-11])
test_sc_pf2 <- pf2_transform(df_test, target_name = "target")
test_sc_pf2$target = df_test$target
test_sc_pf2 = test_sc_pf2[-20]

```

- run the regression models

```
lr = caret::train(target ~ ., method='lm', data = train_sc_pf2)
train_pred = predict(lr, newdata = train_sc_pf2)
val_pred = predict(lr, newdata = val_sc_pf2)
test_pred = predict(lr, newdata = test_sc_pf2)
```

- Evaluate the results

```
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[, sapply(res, is.numeric)] <- round(res[, sapply(res, is.numeric)], 4)
knitr::kable(res)
```

Data	RMSE
Train	44.95
Validation	71.05
Test	77.25

## Parameter Shrinkage with Ridge Regression

```
ridge <- caret::train(y = train_sc_pf2$target, x = train_sc_pf2[-10],
  method = 'glmnet',
  tuneGrid = expand.grid(alpha = 0, lambda = 1)
)
train_pred = predict(ridge, newdata = train_sc_pf2)
val_pred = predict(ridge, newdata = val_sc_pf2)
test_pred = predict(ridge, newdata = test_sc_pf2)

res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
```

```
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	RMSE
Train	46.66
Validation	59.19
Test	57.10

- Best lambda

hyper-parameter tuning

```
parameters <- c(seq(0.1, 2, by =0.1) , seq(2, 5, 0.5) , seq(5, 25, 1))

ridge<-caret::train(y = train_sc_pf2$target,
                    x = train_sc_pf2[-10],
                    method = 'glmnet',
                    tuneGrid = expand.grid(alpha = 0, lambda = parameters) ,
                    metric = "Rsquared"
                )

lasso<-caret::train(y = train_sc_pf2$target,
                    x = train_sc_pf2[-10],
                    method = 'glmnet',
                    tuneGrid = expand.grid(alpha = 1, lambda = parameters) ,
                    metric = "Rsquared"
                )

paste(" Ridge Best lambda = ",ridge$finalModel$lambdaOpt)
```

```
## [1] " Ridge Best lambda = 25"
```

```
paste(" Lasso Best lambda = ",lasso$finalModel$lambdaOpt)
```

```
## [1] " Lasso Best lambda = 3.5"
```

## Run best ridge regression

```
ridge_best<-caret::train(y = train_sc_pf2$target,
                        x = train_sc_pf2[-10],
                        method = 'glmnet',
                        tuneGrid = expand.grid(alpha = 0, lambda = ridge$finalModel$lambdaOpt))
train_pred = predict(ridge_best,newdata = train_sc_pf2)
val_pred = predict(ridge_best,newdata = val_sc_pf2)
test_pred = predict(ridge_best,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
```

```

y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	49.03
Validation	57.48
Test	55.76

```

res = data.frame()
y = R2(train_pred, df_train$target)
res = rbind(res, c("Train", y))
y = R2(val_pred, df_val$target)
res = rbind(res, c("Validation", y))
y = R2(test_pred, df_test$target)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	R_2
Train	0.6275
Validation	0.3794
Test	0.4754

## Lasso Regression

```

lasso<-caret::train(y = train_sc_pf2$target,
  x = train_sc_pf2[-10],
  method = 'glmnet',
  tuneGrid = expand.grid(alpha = 1, lambda = 0))
train_pred = predict(lasso, newdata = train_sc_pf2)
val_pred = predict(lasso, newdata = val_sc_pf2)
test_pred = predict(lasso, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))

```



```

y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[, sapply(res, is.numeric)] <- round(res[, sapply(res, is.numeric)], 4)
knitr::kable(res)

```

Data	RMSE
Train	45.42
Validation	62.98
Test	65.22

## Run best Lasso regression

- Find the best hyper parameter lambda

```

parameters <- c(seq(0.1, 2, by = 0.1), seq(2, 5, 0.5), seq(5, 25, 1))
lasso <- caret::train(y = train_sc_pf2$target,
                      x = train_sc_pf2[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 1, lambda = parameters),
                      metric = "Rsquared"
)
paste(" Lasso Best lambda = ", lasso$finalModel$lambdaOpt)

```

```
## [1] " Lasso Best lambda = 6"
```

- Run the lasso model with the tuned lambda parameter

```

lasso_best <- caret::train(y = train_sc_pf2$target,
                          x = train_sc_pf2[-10],
                          method = 'glmnet',
                          tuneGrid = expand.grid(alpha = 1, lambda = lasso$finalModel$lambdaOpt))
train_pred = predict(lasso_best, newdata = train_sc_pf2)
val_pred = predict(lasso_best, newdata = val_sc_pf2)
test_pred = predict(lasso_best, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type

```

```
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	RMSE
Train	53.35
Validation	57.76
Test	53.54

- R2 with best lasso

```
res = data.frame()
y = R2(train_pred,df_train$target)
res = rbind(res,c("Train",y))
y = R2(val_pred,df_val$target)
res = rbind(res,c("Validation",y))
y = R2(test_pred,df_test$target)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","R_2")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)
```

Data	R_2
Train	0.5569
Validation	0.3648
Test	0.5073

## Lasso non zero coefficients

```
df = data.frame(
  lasso = as.data.frame.matrix(coef(lasso$finalModel, lasso$finalModel$lambdaOpt))
)
df =subset(df, s1>0.1)
df$var = row.names(df)
knitr::kable(df[order(-df$s1),c(2,1)][2])
```

	s1
(Intercept)	153.5906
bp	24.8924
s6	19.8741
s1	12.1702
bmi.s6	3.3431

	s1
sex.s3	1.0124
age.s6	0.2944

## Neural Network

- Large network

```
mlp <- caret::train(target ~ ., data = train_sc_pf2, method='mlp', size=1000)

train_pred = predict(mlp, newdata = train_sc_pf2)
val_pred = predict(mlp, newdata = val_sc_pf2)
test_pred = predict(mlp, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
res = rbind(res, c("Test", y))
colnames(res) = c("Data", "RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[, sapply(res, is.numeric)] <- round(res[, sapply(res, is.numeric)], 4)
knitr::kable(res)
```

Data	RMSE
Train	93.28
Validation	91.27
Test	94.74

- Smaller network

```
mlp <- caret::train(target ~ ., data = train_sc_pf2, method='mlp', size=50)
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.
```

```
train_pred = predict(mlp, newdata = train_sc_pf2)
val_pred = predict(mlp, newdata = val_sc_pf2)
test_pred = predict(mlp, newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target, train_pred)
res = rbind(res, c("Train", y))
y = rmse(df_val$target, val_pred)
res = rbind(res, c("Validation", y))
y = rmse(df_test$target, test_pred)
```

```

res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	87.47
Validation	77.94
Test	83.31

## Regression Trees

```

dtr <- caret::train(target ~ .,data = train_sc_pf2,method='rpart')

train_pred = predict(dtr,newdata = train_sc_pf2)
val_pred = predict(dtr,newdata = val_sc_pf2)
test_pred = predict(dtr,newdata = test_sc_pf2)
res = data.frame()
y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	59.71
Validation	66.82
Test	60.15

- with a maximum depth of 2

```

dtr <- caret::train(target ~ .,data = train_sc_pf2,method='rpart',
                    control = list(max_depth=2))

train_pred = predict(dtr,newdata = train_sc_pf2)
val_pred = predict(dtr,newdata = val_sc_pf2)
test_pred = predict(dtr,newdata = test_sc_pf2)
res = data.frame()

```

```

y = rmse(df_train$target,train_pred)
res = rbind(res,c("Train",y))
y = rmse(df_val$target,val_pred)
res = rbind(res,c("Validation",y))
y = rmse(df_test$target,test_pred)
res = rbind(res,c("Test",y))
colnames(res) = c("Data","RMSE")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Data	RMSE
Train	59.71
Validation	66.82
Test	60.15

## Harder problem

- Read data

```
df <- read.csv("../data/ENB2012_data.csv")
```

- Data preparation

```

dataset = df[-c(11,12)]
N = nrow(dataset)
cut1 = floor(0.6*N)
cut2 = floor(0.8*N)
index = sample(1:N)
train_index = index[1:cut1]
val_index = index[(cut1+1):cut2]
test_index = index[(cut2+1):N]
df_train = dataset[train_index,]
df_val = dataset[val_index,]
df_test = dataset[test_index,]
sapply(list(df_train,df_val,df_test),nrow)

```

```
## [1] 460 154 154
```

```
df_train[-10] = scale(df_train[-10])
```

- Build models

```

res = data.frame()
# Linear Regression
lr = caret::train(Y2 ~ ., method='lm', data = df_train)
train_pred = predict(lr, newdata = df_train)
val_pred = predict(lr, newdata = scale(df_val[-10]))
test_pred = predict(lr, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Linear Regression", y1, y2, y3))
# Ridge Regression
ridge <- caret::train(y = df_train$Y2, x = df_train[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 0, lambda = 1)
)
train_pred = predict(ridge, newdata = df_train)
val_pred = predict(ridge, newdata = scale(df_val[-10]))
test_pred = predict(ridge, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Ridge Regression", y1, y2, y3))
# Lasso Regression
lasso <- caret::train(y = df_train$Y2, x = df_train[-10],
                      method = 'glmnet',
                      tuneGrid = expand.grid(alpha = 1, lambda = 0)
)
train_pred = predict(lasso, newdata = df_train)
val_pred = predict(lasso, newdata = scale(df_val[-10]))
test_pred = predict(lasso, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Lasso Regression", y1, y2, y3))
# Neural Net
nn <- caret::train(y = df_train$Y2, x = df_train[-10],
                   method = 'mlp')
train_pred = predict(nn, newdata = df_train)
val_pred = predict(nn, newdata = scale(df_val[-10]))
test_pred = predict(nn, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)
res = rbind(res, c("Neural Net", y1, y2, y3))
# Regression Tree
dt <- caret::train(y = df_train$Y2, x = df_train[-10],
                   method = 'rpart')
train_pred = predict(dt, newdata = df_train)
val_pred = predict(dt, newdata = scale(df_val[-10]))
test_pred = predict(dt, newdata = scale(df_test[-10]))
y1 = rmse(df_train$Y2, train_pred)
y2 = rmse(df_val$Y2, val_pred)
y3 = rmse(df_test$Y2, test_pred)

```

```

res = rbind(res,c("Regression Tree",y1,y2,y3))

colnames(res) = c("Model","Train","Validation","Test")
# Following converts appropriate columns to numeric type
res[-1] = lapply(res[-1], FUN = function(y){as.numeric(y)})
# Following rounds numeric values to 4 digits
res[,sapply(res, is.numeric)] <-round(res[,sapply(res, is.numeric)],4)
knitr::kable(res)

```

Model	Train	Validation	Test
Linear Regression	2.007	1.829	1.995
Ridge Regression	2.368	2.076	2.256
Lasso Regression	2.008	1.826	1.992
Neural Net	3.855	4.002	3.771
Regression Tree	3.164	2.771	3.100